

In [1]:

```
1 import keras
2 from keras.datasets import mnist
3 from keras.models import Sequential
4 from keras.layers import Dense, Dropout, Flatten
5 from keras.layers import Conv2D, MaxPooling2D
6 from keras.layers import Input, Dense
7 from keras.optimizers import RMSprop
8 from keras import backend as K
9 from keras.models import Model
10 from keras.datasets import mnist
11 from keras import regularizers
12 from keras.utils import np_utils
13 #from keras.utils import to_categorical
14 import tensorflow.python as tf
15 from tensorflow import keras
16
17 from sklearn.linear_model import LogisticRegression
18 from sklearn.ensemble import RandomForestClassifier
19 from sklearn.svm import LinearSVC
20 from sklearn.metrics import accuracy_score
21
22 import numpy as np
23 import matplotlib.pyplot as plt
24
25 # def plot_traincurve(history):
26 #     colors = {'loss':'r', 'acc':'b', 'val_loss':'m', 'val_acc':'g'}
27 #     plt.figure(figsize=(10,6))
28 #     plt.title("Training Curve")
29 #     plt.xlabel("Epoch")
30
31 #     for measure in history.keys():
32 #         color = colors[measure]
33 #         ln = len(history[measure])
34 #         plt.plot(range(1,ln+1), history[measure], color + '-', label=measure)
35
36 #     plt.legend(loc='upper left', scatterpoints = 1, frameon=False)
```

In []:

1

Neural network original data

In [2]:

```

1  batch_size = 128
2  num_classes = 10
3  epochs = 30
4
5  # input image dimensions
6  img_rows, img_cols = 28, 28
7
8  # the data, split between train and test sets
9  (x_train, y_train), (x_test, y_test) = mnist.load_data()
10
11 x_train = x_train.reshape(60000, 784)
12 x_test = x_test.reshape(10000, 784)
13 x_train = x_train.astype('float32')
14 x_test = x_test.astype('float32')
15 x_train /= 255
16 x_test /= 255
17 print(x_train.shape[0], 'train samples')
18 print(x_test.shape[0], 'test samples')
19
20 # convert class vectors to binary class matrices
21 y_train = keras.utils.to_categorical(y_train, num_classes)
22 y_test = keras.utils.to_categorical(y_test, num_classes)
23
24 model = Sequential()
25 model.add(Dense(512, activation='relu', input_shape=(784,)))
26 model.add(Dropout(0.2))
27 model.add(Dense(512, activation='relu'))
28 model.add(Dropout(0.2))
29 model.add(Dense(num_classes, activation='softmax'))
30
31 model.summary()
32
33 model.compile(loss='categorical_crossentropy',
34               optimizer=RMSprop(),
35               metrics=['accuracy'])
36
37 model.fit(x_train, y_train,
38         batch_size=batch_size,
39         epochs=epochs,
40         verbose=1,
41         validation_data=(x_test, y_test))
42 score = model.evaluate(x_test, y_test, verbose=0)
43
44 print('Test loss:', score[0])
45 print('Test accuracy:', score[1])
46 # plot_traincurve(history)

```

60000 train samples

10000 test samples

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 512)	401920
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262656

dropout_1 (Dropout)	(None, 512)	0
---------------------	-------------	---

dense_2 (Dense)	(None, 10)	5130
-----------------	------------	------

=====
Total params: 669,706

Trainable params: 669,706

Non-trainable params: 0

Epoch 1/30

469/469 [=====] - 21s 21ms/step - loss: 0.434

6 - accuracy: 0.8603 - val_loss: 0.1077 - val_accuracy: 0.9679

Epoch 2/30

469/469 [=====] - 8s 18ms/step - loss: 0.1086

- accuracy: 0.9666 - val_loss: 0.0755 - val_accuracy: 0.9776

Epoch 3/30

469/469 [=====] - 10s 20ms/step - loss: 0.075

1 - accuracy: 0.9772 - val_loss: 0.0640 - val_accuracy: 0.9819

Epoch 4/30

469/469 [=====] - 8s 17ms/step - loss: 0.0575

- accuracy: 0.9825 - val_loss: 0.0695 - val_accuracy: 0.9811

Epoch 5/30

469/469 [=====] - 8s 16ms/step - loss: 0.0478

- accuracy: 0.9853 - val_loss: 0.0697 - val_accuracy: 0.9816

Epoch 6/30

469/469 [=====] - 9s 19ms/step - loss: 0.0405

- accuracy: 0.9874 - val_loss: 0.0764 - val_accuracy: 0.9810

Epoch 7/30

469/469 [=====] - 8s 18ms/step - loss: 0.0339

- accuracy: 0.9890 - val_loss: 0.0712 - val_accuracy: 0.9816

Epoch 8/30

469/469 [=====] - 8s 16ms/step - loss: 0.0286

- accuracy: 0.9909 - val_loss: 0.0879 - val_accuracy: 0.9830

Epoch 9/30

469/469 [=====] - 7s 15ms/step - loss: 0.0271

- accuracy: 0.9921 - val_loss: 0.0775 - val_accuracy: 0.9836

Epoch 10/30

469/469 [=====] - 8s 16ms/step - loss: 0.0259

- accuracy: 0.9925 - val_loss: 0.0970 - val_accuracy: 0.9810

Epoch 11/30

469/469 [=====] - 8s 16ms/step - loss: 0.0234

- accuracy: 0.9932 - val_loss: 0.1008 - val_accuracy: 0.9816

Epoch 12/30

469/469 [=====] - 8s 16ms/step - loss: 0.0244

- accuracy: 0.9925 - val_loss: 0.1068 - val_accuracy: 0.9817

Epoch 13/30

469/469 [=====] - 8s 16ms/step - loss: 0.0211

- accuracy: 0.9937 - val_loss: 0.0997 - val_accuracy: 0.9850

Epoch 14/30

469/469 [=====] - 7s 16ms/step - loss: 0.0214

- accuracy: 0.9942 - val_loss: 0.1001 - val_accuracy: 0.9828

Epoch 15/30

469/469 [=====] - 8s 16ms/step - loss: 0.0230

- accuracy: 0.9943 - val_loss: 0.1107 - val_accuracy: 0.9822

Epoch 16/30

469/469 [=====] - 7s 15ms/step - loss: 0.0181

- accuracy: 0.9944 - val_loss: 0.1116 - val_accuracy: 0.9834

Epoch 17/30

469/469 [=====] - 7s 16ms/step - loss: 0.0174

- accuracy: 0.9951 - val_loss: 0.1173 - val_accuracy: 0.9835

Epoch 18/30

469/469 [=====] - 7s 15ms/step - loss: 0.0157

```
- accuracy: 0.9951 - val_loss: 0.1094 - val_accuracy: 0.9827
Epoch 19/30
469/469 [=====] - 7s 16ms/step - loss: 0.0143
- accuracy: 0.9957 - val_loss: 0.1108 - val_accuracy: 0.9845
Epoch 20/30
469/469 [=====] - 8s 17ms/step - loss: 0.0166
- accuracy: 0.9955 - val_loss: 0.1281 - val_accuracy: 0.9826
Epoch 21/30
469/469 [=====] - 9s 20ms/step - loss: 0.0136
- accuracy: 0.9962 - val_loss: 0.1291 - val_accuracy: 0.9838
Epoch 22/30
469/469 [=====] - 8s 16ms/step - loss: 0.0144
- accuracy: 0.9962 - val_loss: 0.1226 - val_accuracy: 0.9846
Epoch 23/30
469/469 [=====] - 7s 15ms/step - loss: 0.0145
- accuracy: 0.9963 - val_loss: 0.1302 - val_accuracy: 0.9825
Epoch 24/30
469/469 [=====] - 7s 15ms/step - loss: 0.0133
- accuracy: 0.9964 - val_loss: 0.1469 - val_accuracy: 0.9837
Epoch 25/30
469/469 [=====] - 8s 16ms/step - loss: 0.0151
- accuracy: 0.9962 - val_loss: 0.1532 - val_accuracy: 0.9803
Epoch 26/30
469/469 [=====] - 7s 16ms/step - loss: 0.0118
- accuracy: 0.9970 - val_loss: 0.1461 - val_accuracy: 0.9838
Epoch 27/30
469/469 [=====] - 9s 20ms/step - loss: 0.0122
- accuracy: 0.9969 - val_loss: 0.1504 - val_accuracy: 0.9849
Epoch 28/30
469/469 [=====] - 8s 16ms/step - loss: 0.0160
- accuracy: 0.9961 - val_loss: 0.1439 - val_accuracy: 0.9860
Epoch 29/30
469/469 [=====] - 7s 15ms/step - loss: 0.0131
- accuracy: 0.9970 - val_loss: 0.1637 - val_accuracy: 0.9820
Epoch 30/30
469/469 [=====] - 7s 16ms/step - loss: 0.0125
- accuracy: 0.9971 - val_loss: 0.1537 - val_accuracy: 0.9838
Test loss: 0.15369866788387299
Test accuracy: 0.9837999939918518
```

Autoencoder

In [3]:

```

1  # this is the size of our encoded representations
2  encoding_dim = 32
3
4  # this is our input placeholder
5  input_img = Input(shape=(784,))
6  # "encoded" is the encoded representation of the input
7  encoded = Dense(encoding_dim, activation='relu')(input_img)
8  # "decoded" is the lossy reconstruction of the input
9  decoded = Dense(784, activation='sigmoid')(encoded)
10
11 # this model maps an input to its reconstruction
12 autoencoder = Model(input_img, decoded)
13
14 # this model maps an input to its encoded representation
15 encoder = Model(input_img, encoded)
16
17 # create a placeholder for an encoded (32-dimensional) input
18 encoded_input = Input(shape=(encoding_dim,))
19 # retrieve the last layer of the autoencoder model
20 decoder_layer = autoencoder.layers[-1]
21 # create the decoder model
22 decoder = Model(encoded_input, decoder_layer(encoded_input))
23
24 autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
25
26 autoencoder.fit(x_train, x_train,
27                 epochs=30,
28                 batch_size=256,
29                 shuffle=True,
30                 validation_data=(x_test, x_test))
31 # plot_traincurve(history.history)
32
33 encoded_imgs_train = encoder.predict(x_train)
34 encoded_imgs_test = encoder.predict(x_test)
35 decoded_imgs = decoder.predict(encoded_imgs_test)

```

Epoch 1/30

235/235 [=====] - 3s 10ms/step - loss: 0.3825

- val_loss: 0.1892

Epoch 2/30

235/235 [=====] - 2s 8ms/step - loss: 0.1793

- val_loss: 0.1521

Epoch 3/30

235/235 [=====] - 2s 9ms/step - loss: 0.1478

- val_loss: 0.1328

Epoch 4/30

235/235 [=====] - 2s 9ms/step - loss: 0.1310

- val_loss: 0.1210

Epoch 5/30

235/235 [=====] - 2s 8ms/step - loss: 0.1203

- val_loss: 0.1133

Epoch 6/30

235/235 [=====] - 2s 9ms/step - loss: 0.1134

- val_loss: 0.1075

Epoch 7/30

235/235 [=====] - 2s 10ms/step - loss: 0.1075

- val_loss: 0.1033

Epoch 8/30

235/235 [=====] - 2s 8ms/step - loss: 0.1036

```
- val_loss: 0.1001
Epoch 9/30
235/235 [=====] - 2s 9ms/step - loss: 0.1008
- val_loss: 0.0978
Epoch 10/30
235/235 [=====] - 2s 9ms/step - loss: 0.0985
- val_loss: 0.0961
Epoch 11/30
235/235 [=====] - 2s 9ms/step - loss: 0.0971
- val_loss: 0.0949
Epoch 12/30
235/235 [=====] - 2s 10ms/step - loss: 0.0961
- val_loss: 0.0942
Epoch 13/30
235/235 [=====] - 2s 10ms/step - loss: 0.0952
- val_loss: 0.0936
Epoch 14/30
235/235 [=====] - 2s 9ms/step - loss: 0.0944
- val_loss: 0.0933
Epoch 15/30
235/235 [=====] - 2s 9ms/step - loss: 0.0945
- val_loss: 0.0931
Epoch 16/30
235/235 [=====] - 2s 10ms/step - loss: 0.0940
- val_loss: 0.0928
Epoch 17/30
235/235 [=====] - 2s 10ms/step - loss: 0.0939
- val_loss: 0.0927
Epoch 18/30
235/235 [=====] - 2s 10ms/step - loss: 0.0939
- val_loss: 0.0926
Epoch 19/30
235/235 [=====] - 3s 11ms/step - loss: 0.0937
- val_loss: 0.0925
Epoch 20/30
235/235 [=====] - 2s 11ms/step - loss: 0.0935
- val_loss: 0.0924
Epoch 21/30
235/235 [=====] - 2s 8ms/step - loss: 0.0933
- val_loss: 0.0923
Epoch 22/30
235/235 [=====] - 2s 10ms/step - loss: 0.0933
- val_loss: 0.0923
Epoch 23/30
235/235 [=====] - 2s 9ms/step - loss: 0.0934
- val_loss: 0.0922
Epoch 24/30
235/235 [=====] - 2s 8ms/step - loss: 0.0933
- val_loss: 0.0922
Epoch 25/30
235/235 [=====] - 2s 9ms/step - loss: 0.0933
- val_loss: 0.0921
Epoch 26/30
235/235 [=====] - 2s 9ms/step - loss: 0.0932
- val_loss: 0.0921
Epoch 27/30
235/235 [=====] - 2s 8ms/step - loss: 0.0931
- val_loss: 0.0921
Epoch 28/30
235/235 [=====] - 2s 9ms/step - loss: 0.0930
- val_loss: 0.0921
```

Epoch 29/30

235/235 [=====] - 2s 10ms/step - loss: 0.0932

- val_loss: 0.0920

Epoch 30/30

235/235 [=====] - 2s 10ms/step - loss: 0.0930


- val_loss: 0.0920

In [4]:

```

1 n = 10 # how many digits we will display
2 plt.figure(figsize=(20, 4))
3 for i in range(n):
4     # display original
5     ax = plt.subplot(2, n, i + 1)
6     plt.imshow(x_test[i].reshape(28, 28))
7     plt.gray()
8     ax.get_xaxis().set_visible(False)
9     ax.get_yaxis().set_visible(False)
10
11     # display reconstruction
12     ax = plt.subplot(2, n, i + 1 + n)
13     plt.imshow(decoded_imgs[i].reshape(28, 28))
14     plt.gray()
15     ax.get_xaxis().set_visible(False)
16     ax.get_yaxis().set_visible(False)
17 plt.show()

```



Neural network encoded data

In [5]:

```

1 encoded_imgs_train_normalized = encoded_imgs_train / np.max(encoded_imgs_train)
2 encoded_imgs_test_normalized = encoded_imgs_test / np.max(encoded_imgs_test)
3
4 model = Sequential()
5 model.add(Dense(512, activation='relu', input_shape=(encoding_dim,)))
6 model.add(Dropout(0.2))
7 model.add(Dense(512, activation='relu'))
8 model.add(Dropout(0.2))
9 model.add(Dense(num_classes, activation='softmax'))
10
11 model.summary()
12
13 model.compile(loss='categorical_crossentropy',
14               optimizer=RMSprop(),
15               metrics=['accuracy'])
16
17 model.fit(encoded_imgs_train_normalized, y_train,
18           batch_size=batch_size,
19           epochs=epochs,
20           verbose=1,
21           validation_data=(encoded_imgs_test_normalized, y_test))
22 score = model.evaluate(encoded_imgs_test_normalized, y_test, verbose=0)
23
24 print('Test loss:', score[0])
25 print('Test accuracy:', score[1])
26 # plot_traincurve(history)

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	16896
dropout_2 (Dropout)	(None, 512)	0
dense_6 (Dense)	(None, 512)	262656
dropout_3 (Dropout)	(None, 512)	0
dense_7 (Dense)	(None, 10)	5130

Total params: 284,682
 Trainable params: 284,682
 Non-trainable params: 0

Epoch 1/30
 469/469 [=====] - 6s 11ms/step - loss: 0.9493
 - accuracy: 0.7281 - val_loss: 0.3198 - val_accuracy: 0.9039
 Epoch 2/30
 469/469 [=====] - 5s 10ms/step - loss: 0.3355
 - accuracy: 0.8970 - val_loss: 0.2095 - val_accuracy: 0.9370
 Epoch 3/30
 469/469 [=====] - 4s 9ms/step - loss: 0.2311
 - accuracy: 0.9298 - val_loss: 0.1627 - val_accuracy: 0.9499
 Epoch 4/30
 469/469 [=====] - 5s 10ms/step - loss: 0.1786
 - accuracy: 0.9454 - val_loss: 0.1278 - val_accuracy: 0.9604
 Epoch 5/30
 469/469 [=====] - 5s 12ms/step - loss: 0.1493


```
- accuracy: 0.9530 - val_loss: 0.1512 - val_accuracy: 0.9487
Epoch 6/30
469/469 [=====] - 6s 13ms/step - loss: 0.1328
- accuracy: 0.9589 - val_loss: 0.1029 - val_accuracy: 0.9679
Epoch 7/30
469/469 [=====] - 5s 10ms/step - loss: 0.1159
- accuracy: 0.9638 - val_loss: 0.1160 - val_accuracy: 0.9639
Epoch 8/30
469/469 [=====] - 5s 10ms/step - loss: 0.1085
- accuracy: 0.9654 - val_loss: 0.0915 - val_accuracy: 0.9740
Epoch 9/30
469/469 [=====] - 5s 10ms/step - loss: 0.1017
- accuracy: 0.9686 - val_loss: 0.1083 - val_accuracy: 0.9663
Epoch 10/30
469/469 [=====] - 5s 10ms/step - loss: 0.0955
- accuracy: 0.9700 - val_loss: 0.0851 - val_accuracy: 0.9721
Epoch 11/30
469/469 [=====] - 4s 10ms/step - loss: 0.0896
- accuracy: 0.9716 - val_loss: 0.0776 - val_accuracy: 0.9765
Epoch 12/30
469/469 [=====] - 6s 12ms/step - loss: 0.0848
- accuracy: 0.9734 - val_loss: 0.1003 - val_accuracy: 0.9698
Epoch 13/30
469/469 [=====] - 5s 12ms/step - loss: 0.0822
- accuracy: 0.9739 - val_loss: 0.0828 - val_accuracy: 0.9767
Epoch 14/30
469/469 [=====] - 4s 9ms/step - loss: 0.0816
- accuracy: 0.9747 - val_loss: 0.0891 - val_accuracy: 0.9745
Epoch 15/30
469/469 [=====] - 5s 11ms/step - loss: 0.0784
- accuracy: 0.9755 - val_loss: 0.0928 - val_accuracy: 0.9719
Epoch 16/30
469/469 [=====] - 5s 11ms/step - loss: 0.0754
- accuracy: 0.9758 - val_loss: 0.0738 - val_accuracy: 0.9791
Epoch 17/30
469/469 [=====] - 5s 10ms/step - loss: 0.0694
- accuracy: 0.9781 - val_loss: 0.0701 - val_accuracy: 0.9799
Epoch 18/30
469/469 [=====] - 5s 10ms/step - loss: 0.0707
- accuracy: 0.9786 - val_loss: 0.0759 - val_accuracy: 0.9793
Epoch 19/30
469/469 [=====] - 5s 10ms/step - loss: 0.0697
- accuracy: 0.9782 - val_loss: 0.0762 - val_accuracy: 0.9784
Epoch 20/30
469/469 [=====] - 5s 10ms/step - loss: 0.0674
- accuracy: 0.9799 - val_loss: 0.0782 - val_accuracy: 0.9762
Epoch 21/30
469/469 [=====] - 5s 10ms/step - loss: 0.0666
- accuracy: 0.9795 - val_loss: 0.0845 - val_accuracy: 0.9779
Epoch 22/30
469/469 [=====] - 5s 10ms/step - loss: 0.0643
- accuracy: 0.9790 - val_loss: 0.0859 - val_accuracy: 0.9763
Epoch 23/30
469/469 [=====] - 5s 10ms/step - loss: 0.0608
- accuracy: 0.9807 - val_loss: 0.0818 - val_accuracy: 0.9793
Epoch 24/30
469/469 [=====] - 5s 10ms/step - loss: 0.0610
- accuracy: 0.9810 - val_loss: 0.0754 - val_accuracy: 0.9804
Epoch 25/30
469/469 [=====] - 5s 10ms/step - loss: 0.0585
- accuracy: 0.9818 - val_loss: 0.0695 - val_accuracy: 0.9811
```

```
Epoch 26/30
469/469 [=====] - 5s 10ms/step - loss: 0.0588
- accuracy: 0.9811 - val_loss: 0.0746 - val_accuracy: 0.9790
Epoch 27/30
469/469 [=====] - 5s 10ms/step - loss: 0.0580
- accuracy: 0.9820 - val_loss: 0.0783 - val_accuracy: 0.9788
Epoch 28/30
469/469 [=====] - 5s 10ms/step - loss: 0.0574
- accuracy: 0.9824 - val_loss: 0.0911 - val_accuracy: 0.9767
Epoch 29/30
469/469 [=====] - 4s 9ms/step - loss: 0.0558
- accuracy: 0.9827 - val_loss: 0.0639 - val_accuracy: 0.9831
Epoch 30/30
469/469 [=====] - 5s 11ms/step - loss: 0.0567
- accuracy: 0.9824 - val_loss: 0.0879 - val_accuracy: 0.9780
Test loss: 0.08785980939865112
Test accuracy: 0.9779999852180481
```

SVM original data

In [7]:

```
1 clf_svm = LinearSVC()
2 clf_svm.fit(x_train, np.argmax(y_train, axis=1))
3 y_pred_svm = clf_svm.predict(x_test)
4 acc_svm = accuracy_score(np.argmax(y_test, axis=1), y_pred_svm)
5 print('Linear SVM accuracy: ', acc_svm)
```

Linear SVM accuracy: 0.9181

```
/home/minhvu/anaconda3/lib/python3.8/site-packages/sklearn/svm/_base.p
y:976: ConvergenceWarning: Liblinear failed to converge, increase the
number of iterations.
  warnings.warn("Liblinear failed to converge, increase "
```

SVM encoded data

In [8]:

```
1 clf_svm = LinearSVC()
2 clf_svm.fit(encoded_imgs_train_normalized, np.argmax(y_train, axis=1))
3 y_pred_svm = clf_svm.predict(encoded_imgs_test_normalized)
4 acc_svm = accuracy_score(np.argmax(y_test, axis=1), y_pred_svm)
5 print('Linear SVM accuracy: ', acc_svm)
```

Linear SVM accuracy: 0.8883

Random Forest original data

In [9]:

```
1 clf_rf = RandomForestClassifier()
2 clf_rf.fit(x_train, np.argmax(y_train, axis=1))
3 y_pred_rf = clf_rf.predict(x_test)
4 acc_rf = accuracy_score(np.argmax(y_test, axis=1), y_pred_rf)
5 print('random forest accuracy: ',acc_rf)
```

random forest accuracy: 0.9693

Random Forest encoded data

In [10]:

```
1 clf_rf = RandomForestClassifier()
2 clf_rf.fit(encoded_imgs_train_normalized, np.argmax(y_train, axis=1))
3 y_pred_rf = clf_rf.predict(encoded_imgs_test_normalized)
4 acc_rf = accuracy_score(np.argmax(y_test, axis=1), y_pred_rf)
5 print('random forest accuracy: ',acc_rf)
```

random forest accuracy: 0.9419

In []:

```
1
```